

Temporary Difference: Applications

Saturday, May 14, 2022 9:50 AM

We have now two codes that apply Q table learning using Temporary Difference method for two different Gym environments:

- 1) Mountain car
- 2) Cart Pole

1) Mountain car

The console output looks as follows:

```
e= 213 r_total= -200.0 r_MAX= -200.0 r_prom= -200.0 epsilon= 1
e= 214 r_total= -200.0 r_MAX= -200.0 r_prom= -200.0 epsilon= 1
e= 215 r_total= -200.0 r_MAX= -200.0 r_prom= -200.0 epsilon= 1
e= 216 r_total= -200.0 r_MAX= -200.0 r_prom= -200.0 epsilon= 1
e= 217 r_total= -200.0 r_MAX= -200.0 r_prom= -200.0 epsilon= 1
e= 218 r_total= -200.0 r_MAX= -200.0 r_prom= -200.0 epsilon= 1
e= 219 r_total= -200.0 r_MAX= -200.0 r_prom= -200.0 epsilon= 1
e= 220 r_total= -200.0 r_MAX= -200.0 r_prom= -200.0 epsilon= 1
e= 221 r_total= -200.0 r_MAX= -200.0 r_prom= -200.0 epsilon= 1
```

current episode

episode: the car tries a maximum of 200 actions. If the car arrives the episode is over; if it can't arrive in 200 actions, the episode is over too.

→ The world punishes the agent with -1 for every action taken that does not result in reaching the goal.


r-total prints {

- ↳ If the agent could not reach the goal in 200 steps = 200 punishes → reward = -200
- ↳ If the agent arrives to the goal (less than 200 steps in the episode) = less than 200 punishes = reward greater than -200 i.e. -137, etc.

↳ arriving to the goal means neither punish nor reward → $r += 0$

r-max prints {

- ↳ The maximum reward in all episodes so far: if there was a previous episode with greater r-prom (reward), it appears in following episodes as r-max
- ↳ a reward r-total = -200 in e = # means that in episode # the agent executed 200 actions that did not result in the goal.



r-prom = -200 test show no learning {

- ↳ r-prom: every 100 episodes, we test the agent in 20 episodes and compute the average of those 20 tests every 100th episode.

↳ we aim for greater reward prints every time.

→ How is the agent deciding the actions?

table Q(s,a) initially with random numbers [0,1]
self.q = np.random.rand(self.Nx+1, self.Nv+1, self.Na)

↳ The agent thus decides what to do with Q position (max)
Since initially with randoms in Q → random actions in the beginning
In agent:

```
def action(self, s, env):
    return self.q.action(s) → the decision of action is Q
                             which computes the max()
```

As learning progresses, the agent will start taking decisions with informed actions: the random Q values get adjusted better.

When the agent is LEARNING, the agent can take actions in TWO WAYS:


→ Exploration: When the agent does not know and takes random actions, the agent is exploring
 the agent, when taking random actions, is exploring the state space.
 → Moves with no preference
 → We don't use Q really, we use random actions.

→ Exploitation: When the agent is taking actions based on the knowledge it has (Q table), we call it exploitation.
 → Once the agent knows the way to achieve max reward (Optimal policy), it is time to take that path: exploitation.
 → We use Q which tells us what to do.

In the program, the agent always uses Q, but in the beginning is random → exploration. As learning goes on, exploration starts to shift to exploitation.

↳ We can thus define in functional form how much exploration and exploitation the agent does.

We define a probability distribution for those two concepts:



probability of exploration + probability of exploitation = 1

always: $P(\text{Exploration} | t) + P(\text{Exploitation} | t) = 1$

all with respect to a particular transition (conditioned to transition)

=

$p(\text{exploration}) + p(\text{exploitation}) = 1$
 at point t at point t

→ Exploration is how the agent takes an action

→ Say $P(\text{Expl}) = 0.8$ and $P(\text{Expl}) = 0.2$ that means
 when agent wants a new action: 0.8 prob it is from Q
 0.2 prob it is random
 take a number: ← on that action taken
 $n > 0.8 \rightarrow \text{Rand}$
 $n < 0.8 \rightarrow Q$
 { 80% from Q
 20% from random

→ We don't know Q's real functional form ~~if~~
 a way to solve this is to define $P(\text{expl})$ and $P(\text{explt})$ curves.

→ In this example, 0 to 2 transition have $P(\text{expl}) = 1$

3 to 6 this prob changes


→ $P(\text{explr})$ does not reach 0, it reaches ϵ_{MIN} .

ϵ_{MAX} is maximum probability, or 1.

- $P(\text{explr})$ reaches ϵ_{MIN}
 - $P(\text{explt})$ reaches ϵ_{MAX}
- because we want that when the agent learned, agent still takes random sometimes to adjust it: since agent very likely didn't explore all space and thus Q can still be adjusted a bit.

there may be other path that is better along the way.

→ What functional form should $P(\text{explt})$ and $P(\text{explr})$ have?



→ if we always choose random action, the agent will never learn the goal. It might reach it, but never learn (Q-learning)

→ all this refers to Q but can also ask V .

When program finishes

`e= 9999 r_total= -200.0 r_MAX= -115.0 r_prom= -199.65 epsilon= 1`

this episode never won

Some episodes were won


tests: learned a little

You can increase the MAX-NUM-EP to 20,000 to improve this

After this we can test one episode: some will arrive some won't

→ One policy only, but initial conditions change (gym issue)

$p(\text{explr}) = 0.2$ $p(\text{explt}) = 0.8$ at current t



— If we never explore and always exploit, Q matrix will move it with a simple preference (not optimal) and with exploration, the Q gets updated with rewards and it spreads across the matrix like a ripple = agent finds optimal

↳ when agent gets reward/punish, Q gets updated like a ripple